

# Artificial Intelligence

Alex Day

September 24, 2020

## Contents

<b>1</b>	<b>AI Basics</b>	<b>2</b>
1.1	Agents . . . . .	2
1.2	Searching . . . . .	3
1.3	State Space Graphs vs Search Trees . . . . .	3
<b>2</b>	<b>Uninformed Search</b>	<b>4</b>
2.1	Breadth First Search . . . . .	4
2.2	Depth First Search . . . . .	4
2.3	Iterative Deepening . . . . .	4
2.4	Uniform Cost Search . . . . .	4
2.5	Search Algorithm Evaluation . . . . .	5
2.5.1	BFS Properties . . . . .	5
2.5.2	DFS Properties . . . . .	6
2.5.3	UCS properties . . . . .	6
<b>3</b>	<b>Informed Search</b>	<b>6</b>
3.1	Search Heuristic . . . . .	6
3.2	Greedy Search . . . . .	7
3.3	A* . . . . .	7
<b>4</b>	<b>Probabilities</b>	<b>8</b>
<b>5</b>	<b>Markov Decision Process</b>	<b>8</b>
5.1	Stochasticity . . . . .	8
5.2	Gridworld . . . . .	8
5.3	Stochastic motion model . . . . .	9
5.4	Simple Game . . . . .	10
5.5	Markov Decision Process . . . . .	10

5.6	Policy . . . . .	11
5.7	Discounting . . . . .	12
5.8	Avoiding Infinite Rewards . . . . .	12
5.9	Revisiting MDPs . . . . .	12
5.10	Solving MDPs . . . . .	13
	5.10.1 Value Iteration . . . . .	13
	5.10.2 Policy Iteration . . . . .	14
<b>6</b>	<b>Reinforcement Learning</b>	<b>15</b>
6.1	Definition . . . . .	15
6.2	Framework . . . . .	15
6.3	Model-Based Learning . . . . .	16
	6.3.1 Model-Based Monte Carlo . . . . .	16
6.4	Model-Free Learning . . . . .	16
	6.4.1 Example . . . . .	17
6.5	Q-Learning . . . . .	17
	6.5.1 Active RL . . . . .	17
	6.5.2 Tabular Q-Learning . . . . .	17
	6.5.3 Properties . . . . .	17
	6.5.4 Exploration vs Exploitation . . . . .	18
<b>7</b>	<b>Adversarial Games</b>	<b>18</b>
7.1	Types of Games . . . . .	18
7.2	Deterministic Games . . . . .	19
7.3	Why are games hard? . . . . .	19
7.4	Zero Sum Games . . . . .	19
7.5	General Games . . . . .	19

# 1 AI Basics

## 1.1 Agents

- An **agent** is an entity that perceives and acts
- A **rational agent** selects actions that achieve the best (expected) outcome
- **Reflex agents** consider how the world is but do **not** consider future consequences of their actions
  - Can sometimes be rational, although not always

- **Planning agents** consider how the world would be based upon their actions and have some goal
  - Decisions are based on hypothesized consequences of actions
  - Not always the **best** action so they're not always rational

## 1.2 Searching

- In a Discrete Search Problem we are given:
  - A finite state space
  - A finite action space
  - A cost function
    - \*  $Cost = C(Action, State, FutureState)$
    - \* The cost of an action is defined as the cost of moving from a state to some future state through that action
  - A transition model
    - \*  $FutureState = Transition(Action, CurrentState)$
  - Start state and a goal state or goal test
  - We seek to find a minimum cost solution: a sequence of actions that lead from the start to the goal
  - We assume the cost of the solution is equal to the sum of the cost of each step

## 1.3 State Space Graphs vs Search Trees

- State Space Graphs
  - The state space forms a directed graph where the nodes are states and the edges are actions
  - Each state occurs only once
  - Goal test is a set of nodes
  - Rarely can build it in memory
- Search Trees
  - Root has the start state
  - Branches are actions

- The nodes show states but correspond to local PLANS
- Search trees can be expanded until the solution is found
  - \* Leaf nodes are called the frontier or the open list
  - \* Leaf nodes are nodes that have unexplored options

## **2 Uninformed Search**

### **2.1 Breadth First Search**

- Expand shallowest node first
- Frontier is a FIFO queue

### **2.2 Depth First Search**

- Always expand the deepest node first
- Frontier is a LIFO queue (stack)

### **2.3 Iterative Deepening**


- Run DFS with depth limit 1
- Run DFS with depth limit 2
- Run DFS with depth limit ...
- DFS space complexity with BFS time complexity

### **2.4 Uniform Cost Search**

- Expand least-cost node first
- Frontier is a priority queue
- Issues
  - Explores in all directions
  - No goal-oriented expansion

## 2.5 Search Algorithm Evaluation

- Completeness - Does this always find a solution if one exists
- Optimal - Find the least cost solution
- Time complexity - Time taken
- Space complexity - Space needed
- Useful quantities
  - $b$  - branching factor of the tree (average number of successors for any node)
  - $m$  - Maximum depth of the state space (meaning there are at max  $b^m$  nodes)
  - $d$  - Depth of the shallowest goal node



### 2.5.1 BFS Properties

- Time complexity  $\approx O(b^d)$
- Space complexity  $\approx O(b^d)$
- It is complete (if  $d$  is finite)
- It is optimal only if step costs are equal or increasing as we move down the tree
- BFS requires a crazy amount of memory and time

### 2.5.2 DFS Properties

- It is complete if  $m$  is finite and the graph is acyclic
- Not optimal
- Time complexity  $O(b^m)$  if  $m \neq \infty$  and terrible if  $m \gg d$
- Space complexity  $O(bm)$

### 2.5.3 UCS properties

- It is optimal
- It is complete if the cost of every action is at least  $\epsilon > 0$
- Time
  - If  $C^*$  is the optimal cost the effective depth is  $\frac{C^*}{\epsilon}$
  - It takes  $O(b^{\frac{C^*}{\epsilon}})$  time and space

## 3 Informed Search

- **Informed Search Methods** use problem specific knowledge to solve a problem better
- **Idea:** Use an *evaluation* function  $f(n)$  for each node  $n$ 
  - Estimate “desirability” of each node
- Open is a priority queue sorted by increasing  $f$ -cost

### 3.1 Search Heuristic

- A heuristic function  $h(n)$ 
  - Estimates how close the state at node  $n$  is to the goal state
  - Designed for a particular search problem
  - Common heuristics: Manhattan distance, Euclidean distance, etc.

### 3.2 Greedy Search

- Expand the node that appears to be closest to the goal at each step
- $f(n) = h(n)$
- Complete
- Not Optimal
- Time -  $O(b^m)$
- Space -  $O(b^m)$

### 3.3 A\*

- Guide the search while avoid expanding expensive paths
- Evaluation function  $f(n) = g(n) + h(n)$
- Admissible heuristics
  - Never overestimate true cost of the goal
- Consistent heuristics
  - $h(n) \leq c(n, a, n') + h(n')$
  - Where  $c$  is a step cost function
  - All consistent heuristics are admissible
- Most of the work in A\* lies on finding admissible heuristics
  - We can often find these by solving a *relaxed* version of the problem
    - \* The **key** idea is the optimal solution cost of the relaxed problem is no greater than the optimal solution cost of the real problem
- Given two heuristics  $h_1$  and  $h_2$  if  $h_2(n) \geq h_1(n) \forall n$  then  $h_2$  **dominates**  $h_1$  and is better for search
- Given  $m$  admissible heuristics  $h_1, h_2, \dots, h_m$  then  $h(n) = \max(h_1(n), h_2(n), \dots, h_m(n))$  is also admissible and dominates any  $h_i$
- A\* has extensions that allow incremental, anytime, and pruning approaches

## 4 Probabilities

- A random variable  $X$ , represents an event whose outcome is unknown
- $P$  is a probability distribution that assigns weight to outcomes
  - $X = \text{weather tomorrow}$
  - $X \in \{\text{sunshine}, \text{rain}, \text{thunder}\}$
  - $P(X = \text{sunshine}) = 0.5, P(X = \text{rain}) = 0.25, P(X = \text{thunder}) = 0.25$
- Probabilities are non-neg and sum to one
- As we get more evidence the probabilities may change
- The **expected value**,  $E$ , of a function of a random variable is the average, weighted by the probability distribution over outcomes

## 5 Markov Decision Process


### 5.1 Stochasticity

- Sometimes you cannot rely that a given action from a specific state may always take you to a certain state
- How can we act optimally in the face of randomness

### 5.2 Gridworld

- Noisy motion model
  - 80% the action N takes the agent North (if there is no wall)
  - 10% N goes West, 10% N goes east
  - If there is a wall the agent stays put
- The agent receives rewards at each step and a big reward if it exits at +1 and a bad reward if it exits at -1






### 5.3 Stochastic motion model




Stochastic world



## 5.4 Simple Game

- At each round:
  1. Stay or quit
  2. If quit: you get \$10
  3. If stay: you get \$5 and then roll a die
    - (a) If the result is 1 or 2 the game ended
    - (b) Otherwise the game continues to the next round




In a Markov Decision Process there are state nodes (in blue), chance nodes (in green), choice edges (in black text), and reward edges with the probability and reward (in red)

## 5.5 Markov Decision Process

- A set of states  $S$
- A set of actions  $A$

- A transition function  $T(s, a, s')$ 
  - Also called the model or dynamics
  - Sometimes  $P(s'|s, a)$
- A reward function  $R(s, a, s')$ 
  - Sometimes just  $R(s)$  or  $R(s')$
  - A start state  $s_0$  (and maybe a terminal state)
- MDPs are non-deterministic search problems



- In a MDP, “Markov” means that the action outcomes depend only on the current state

$$\begin{aligned}
 &P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0) \\
 &= \\
 &P(S_{t+1} = s' | S_t = s_t, A_t = a_t)
 \end{aligned}$$

- This is just like search: in the 1st assignment the children of an expanded state depended only on the current node not how you got there

## 5.6 Policy

- In deterministic we seek an optimal sequence of actions from start to goal
- In stochastic we seek an optimal policy  $\pi^* : S \rightarrow A$ 
  - Policy  $\pi$  gives an action for each state
- Following a policy yields a random path

- The utility,  $U$ , of a policy is the (discounted) sum of the rewards along the path
- The goal is to find an optimal policy,  $\pi^*$  that maximizes the expected utility

## 5.7 Discounting

- It's reasonable to maximize the sum of rewards
- It's also reasonable to prefer rewards now to rewards later
- Solution: values of rewards decay exponentially over time based on some discount factor  $0 \leq \gamma \leq 1$
- We discount so that algorithms can converge and have theoretical guarantees

## 5.8 Avoiding Infinite Rewards

- **Problem:** If the game lasts forever, do we get infinite rewards?
- Solutions:
  - Introduce some artificial time horizon  $H$ 
    - \* Gives nonstationary policies ( $\pi$  depends on the time left)
  - Discounting: use  $0 < \gamma < 1$
  - Absorbing state: guarantee that for every policy a terminal state will eventually be reached

## 5.9 Revisiting MDPs

- Same definition as before but with:
  - Discount factor  $\gamma$
  - Horizon  $H$  (can be  $\infty$ )
- MDP quantities so far
  - Policy  $\pi$ : Choice of action for each state
  - Utility  $U_\pi = \sum_{t=0}^H \gamma^t R(s_t)$  Sum of discounted rewards

## 5.10 Solving MDPs

### 5.10.1 Value Iteration


- $\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}[\sum_{t=0}^H \gamma^t R(s_t) | \pi]$

#### 1. Optimal value Function $V^*$

- $V^*(s)$  is the sum of discounted rewards starting in  $s$  and acting optimally

Let's assume:

Noise = 0,  $\gamma = 1$ , living reward = 0



$$V^*(4,3) = 1$$

$$V^*(3,3) = 1$$


$$V^*(2,3) = 1$$

$$V^*(1,1) = 1$$

$$V^*(4,2) = -1$$

Let's assume:

Noise = 0,  $\gamma = 0.9$ , living reward = 0



$$V^*(4,3) = 1$$

$$V^*(3,3) = 0.9$$

$$V^*(2,3) = 0.9 \times 0.9 = 0.81$$


$$V^*(1,1) = 0.9 \times 0.9 \times 0.9 \times 0.9 \times 0.9 = 0.59$$

$$V^*(4,2) = -1$$

- Closed form for  $V^*$  is  $CVV^* = \max_a (R(s, a, s') + \gamma V^*(s'))$

Let's assume:

Noise = 0.2,  $\gamma = 0.9$ , living reward = 0



$$V^*(4,3) = 1$$

$$V^*(3,3) = 0.8 \times 0.9 + 0.1 \times 0.9 \times V^*(3,3) + 0.1 \times 0.9 \times V^*(3,2) \text{ [action East]}$$

$$V^*(2,3) =$$

$$V^*(1,1) =$$

$$V^*(4,2) =$$

2. Bellman Equation  $V^*(s) = \max_a \mathbb{E}[R(s, a, s') + \gamma V^*(s')]$

### 3. Q-Values

- $Q^*(s, a)$  = expected utility starting at state  $s$ , taking action  $a$ , and then acting optimally

• Bellman equation

$$Q^*(s, a) \leftarrow \sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma \max_{a'} Q^*(s', a'))$$

## 5.10.2 Policy Iteration

### 1. Extracting Policy from $V^*$

- It's not obvious
- We need to keep track of the optimal policy during value iteration
- This is called policy extraction, since the policy is informed by the value

### 2. Extracting policy from $Q^*$

- Just take the max from each
- Optimal policy is implicit

### 3. Issues with Value iteration

- Slow  $O(S^2A)$  per iteration
- The “max” at each state rarely changes
- The policy often converges before the actions

### 4. Policy evaluation

- In value iteration we max over all actions to compute optimal values

- If we fixed some policy  $\pi(s)$ , then only one action per state
  - $V^\pi(s)$  = expected total discounted rewards starting in  $s$  and following  $\pi$
  - The value depends now on which policy we fixed
- Iterate and converge at optimal policy rather than optimal value

## 6 Reinforcement Learning

### 6.1 Definition

- MDP but we don't know  $P$  and  $R$

### 6.2 Framework

At every timestep  $t$  the agent picks an action  $a_t$  and the environment changes to a new state  $s_t$  and the agent is given some reward  $r_t$



### 6.3 Model-Based Learning



#### Unknown MDP: Model-Based

Goal	Technique
Compute $V^*, Q^*, \pi^*$	VI/PI on approx. MDP
Evaluate a fixed policy $\pi$	PE on approx. MDP

#### 6.3.1 Model-Based Monte Carlo

1. Learn empirical MDP using Monte Carlo simulation

- (a) Episodic data:  $s_0, a_0, r_1, s_1, a_1, \dots, s_T$
- (b) Estimate transitions and rewards

$$\bullet \hat{P}(s' | s, a) = \frac{\# \text{ times } (s, a, s') \text{ occurs}}{\# \text{ times } (s, a) \text{ occurs}}$$

- (c) Discover each  $\hat{R}$  when we experience  $(s, a, s')$
2. Estimates converge to the truth
  3. Solve using value/policy iteration

### 6.4 Model-Free Learning

Goal	Technique
Compute $V^*, Q^*, \pi^*$	Q-learning
Evaluate a fixed policy $\pi$	MC, TD Learning



### 6.4.1 Example

- 

## 6.5 Q-Learning

- Used to compute  $V^*$ ,  $Q^*$ , and  $\pi^*$  in a model free, unknown MDP

### 6.5.1 Active RL

- Full reinforcement learning
- Learner makes choices
- Exploration vs Exploitation

### 6.5.2 Tabular Q-Learning

- Sample-based Q-value iteration

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} P(s'|s, a)(R(s, a, s') + \gamma \max_{a'} Q_k(s', a'))$$

- Learn Q-values as you go
  - Recieve a sample  $(s, a, r, s')$
  - Consider previous estimate  $\hat{Q}(s, a)$
  - Consider your new sample estimate
    - \*  $\text{target} = R(s, a, s') + \gamma \max_{a'} \hat{Q}(s', a')$
  - Incorporate the new estimate into a running average using some learning rate  $\alpha$ 
    - \*  $\hat{Q}(s, a) \leftarrow (1 - \alpha)\hat{Q}(s, a) + \alpha[\text{target}]$

### 6.5.3 Properties

- Q-learning converges to optimal policy, even if you're acting suboptimally if the policy visits every path
- This is called off-policy learning
  - On policy: Learn values of the policy used to generate samples

- Off policy: Learn the value of another policy
- Caveats
  - Lot of exploration
  - Make the learning rate small enough eventually
  - Basically, in the limit, it doesn't matter how you select actions!

#### 6.5.4 Exploration vs Exploitation

- All states should be visited infinitely often
  - Multi-arm bandit
1. How to sample actions?
    - Choose random actions?
      - You need to visit a lot of cells to get enough data of optimal paths
    - Choose action greedily?
      - Never try the other paths that could be higher value?
    - Answer: start by random and transition to greedy to converge
    - Epsilon-greedy
      - With a (small) probability  $\epsilon$ , act randomly
      - With a (large) probability  $1 - \epsilon$ , act on current policy

## 7 Adversarial Games

### 7.1 Types of Games

- Axes
  - Deterministic vs Stochastic
  - Perfect Information
  - One, two, or more players?
  - Zero sum?



	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information	battleships, blind tictactoe	bridge, poker scrabble

## 7.2 Deterministic Games

- Search Problem
  - States:  $S$  (starts at  $s_0$ )
  - Players:  $P = 1 \dots N$
  - Actions:  $A$
  - Transition Function  $S \times A \rightarrow S$
  - Terminal test:  $S \rightarrow \{\text{true}, \text{false}\}$
  - Terminal utilities:  $S \times P \rightarrow \mathbb{R}$
- Solution for a *player*  $p$  is a policy  $\pi_p : S \rightarrow A$ 
  - $\pi_p$  is defined when it is  $p$ 's turn to play

## 7.3 Why are games hard?

- The utility comes at the end of the game
- Each state is controlled by different players

## 7.4 Zero Sum Games

- Agents have opposite utilities
- Single value one maximizes that the other minimizes
- Adversarial, pure competition

## 7.5 General Games

- Agents have independent utilities
- Cooperation, indifference, competition, etc.